

DATA-236 Sec 12 - Distributed Systems for Data Engineering

HOMEWORK 7

Nandhakumar Apparsamy

018190003

Total Point: 10

GitHub - <https://github.com/Nandha951/DATA-236-HW8-Kafka>

Objective

Demonstrate understanding of Kafka topics, producers, consumers, and partitions by building a simple Node.js + KafkaJS message queue system.

1. Create Two Kafka Topics

Create the following topics:

- o orders-topic
- o payments-topic

Both topics should have:

- o 3 partitions
- o Replication factor of 1 (since it's a single broker setup)

Example:

`docker exec -it kafka-demo_kafka_1 kafka-topics --create --topic orders-topic --partitions 3 --replication-factor 1 --bootstrap-server localhost:9092`

```
(base) spartan@MLK-SCS-P0WGL9N2QF HMB % docker-compose up -d
WARN[0000] /Users/spartan/SJSU/DATA 236/HW Assignment/HW8/docker-compose.yml: the attribute 'version' is obsolete, it will be ignored, please remove it to avoid potential confusion
[+] Running 18/19
  ✓ kafka Pulled
  ✓ zookeeper Pulled
[+] Running 3/3
  ✓ Network hw8_default Created
  ✓ Container hw8-zookeeper-1 Started
  ✓ Container hw8-kafka-1 Started
(base) spartan@MLK-SCS-P0WGL9N2QF HMB % docker exec -it hw8-kafka-1 kafka-topics --create --topic orders-topic --partitions 3 --replication-factor 1 --bootstrap-server localhost:9092
Created topic orders-topic.
(base) spartan@MLK-SCS-P0WGL9N2QF HMB % docker exec -it hw8-kafka-1 kafka-topics --create --topic payments-topic --partitions 3 --replication-factor 1 --bootstrap-server localhost:9092
Created topic payments-topic.
```

2. Develop an Order Producer (orderProducer.js)

Accept CLI input for:

- o orderId

- o orderAmount

Send a JSON message to orders-topic in this structure:

```
{
  "orderId": "string",
  "orderAmount": number,
  "status": "PENDING",
  "timestamp": "ISO8601 formatted date"
}
```

```
const { Kafka } = require('kafkajs')
const { v4: uuidv4 } = require('uuid');

const kafka = new Kafka({
  clientId: 'order-producer',
  brokers: ['localhost:9092']
})

const producer = kafka.producer()

const orderId = process.argv[2] || uuidv4();
const orderAmount = parseFloat(process.argv[3]) || Math.floor(Math.random() * 100) + 1;

const sendMessage = async () => {
  try {
    await producer.connect()
    const message = {
      orderId: orderId,
      orderAmount: orderAmount,
      status: "PENDING",
      timestamp: new Date().toISOString()
    }

    await producer.send({
      topic: 'orders-topic',
      messages: [
        { value: JSON.stringify(message) }
      ],
    })

    console.log(`Sent message: ${JSON.stringify(message)}`)
    await producer.disconnect()
  } catch (error) {
    console.error(error)
  }
}

sendMessage()
```

```
● (base) spartan@MLK-SCS-P8WGL9N2QF HW8 % node orderProducer.js order1 100
{"level":"WARN","timestamp":"2025-03-25T04:44:50.065Z","logger":"kafkajs","message":"KafkaJS v2.0.0 switched de
.js.org/docs/migration-guide-v2.0.0#producer-new-default-partitioner for details. Silence this warning by setti
(node:55368) TimeoutNegativeWarning: -1742877890090 is a negative number.
Timeout duration was set to 1.
(Use `node --trace-warnings ...` to show where the warning was created)
Sent message: {"orderId":"order1","orderAmount":100,"status":"PENDING","timestamp":"2025-03-25T04:44:50.090Z"}
```

```
{"level":"INFO","timestamp":"2025-03-25T04:43:39.251Z","logger":"kafkajs","message":"[Consumer] Starting","groupId":"order-group"}
{"level":"ERROR","timestamp":"2025-03-25T04:43:39.298Z","logger":"kafkajs","message":"[Connection] Response GroupCoordinator(key: 10
{"level":"ERROR","timestamp":"2025-03-25T04:43:39.641Z","logger":"kafkajs","message":"[Connection] Response GroupCoordinator(key: 10
{"level":"INFO","timestamp":"2025-03-25T04:43:43.547Z","logger":"kafkajs","message":"[ConsumerGroup] Consumer has joined the group",
rs-topic":["0,1,2"],"groupProtocol":"RoundRobinAssigner","duration":4287}
{
  topic: 'orders-topic',
  partition: 1,
  offset: '0',
  key: undefined,
  value: {
    orderId: 'order1',
    orderAmount: 100,
    status: 'PENDING',
    timestamp: '2025-03-25T04:44:50.090Z'
  }
}
}
```

3. Develop a Payment Producer (paymentProducer.js)

Accept CLI input for:

- o paymentId
- o orderId
- o paymentAmount

Send a JSON message to payments-topic in this structure:

```
{
  "paymentId": "string",
  "orderId": "string",
  "paymentAmount": number,
  "status": "COMPLETED",
  "timestamp": "ISO8601 formatted date"
}
```

```
const { Kafka } = require('kafkajs')
const { v4: uuidv4 } = require('uuid')

const kafka = new Kafka({
  clientId: 'payment-producer',
  brokers: ['localhost:9092']
})

const producer = kafka.producer()

const paymentId = process.argv[2] || uuidv4()
const orderId = process.argv[3] || uuidv4()
const paymentAmount = parseFloat(process.argv[4]) || Math.floor(Math.random() * 100) + 1

const sendMessage = async () => {
  try {
    await producer.connect()
    const message = {
      paymentId: paymentId,
      orderId: orderId,
      paymentAmount: paymentAmount,
      status: 'COMPLETED',
      timestamp: new Date().toISOString()
    }

    await producer.send({
      topic: 'payments-topic',
      messages: [
        { value: JSON.stringify(message) }
      ],
    })

    console.log(`Sent message: ${JSON.stringify(message)}`)
    await producer.disconnect()
  } catch (error) {
    console.error(error)
  }
}

sendMessage()
```

```
(base) spartan@MLK-SCS-P0WGL9N2QF HW8 % node paymentProducer.js payment1 order1 50
{"level":"WARN","timestamp":"2025-03-25T04:45:01.325Z","logger":"kafkajs","message":"KafkaJS v2.0.0 switched default partitioner. To retain the current behavior, set the environment variable KAFKAJS_PARTITIONER=v1. For more details, see https://kafkajs.github.io/docs/migration-guide-v2.0.0#producer-new-default-partitioner for details. Silence this warning by setting the environment variable KAFKAJS_PARTITIONER=v1."}
(node:55599) TimeoutNegativeWarning: -1742877901343 is a negative number.
Timeout duration was set to 1.
(Use `node --trace-warnings ...` to show where the warning was created)
Sent message: {"paymentId":"payment1","orderId":"order1","paymentAmount":50,"status":"COMPLETED","timestamp":"2025-03-25T04:45:01.343Z"}
```

4. Develop an Order Consumer (orderConsumer.js)

Consume messages from orders-topic

Log the following clearly:

- o Partition number
- o Offset
- o Parsed JSON data fields (orderId, orderAmount, etc.)

```
const { Kafka } = require('kafkajs')

const kafka = new Kafka({
  clientId: 'order-consumer',
  brokers: ['localhost:9092']
})

const consumer = kafka.consumer({ groupId: 'order-group' })

const consume = async () => {
  await consumer.connect()
  await consumer.subscribe({ topic: 'orders-topic', fromBeginning: true })

  await consumer.run({
    eachMessage: async ({ topic, partition, message }) => {
      const parsedMessage = JSON.parse(message.value.toString())
      console.log({
        topic: topic,
        partition: partition,
        offset: message.offset,
        key: message?.key?.toString(),
        value: parsedMessage,
      })
    },
  })
}

consume()
```

```
% (base) spartan@MLK-SCS-P0WGL9N2QF HW8 % node orderConsumer.js
(node:53882) TimeoutNegativeWarning: -1742877819225 is a negative number.
Timeout duration was set to 1.
(Use `node --trace-warnings ...` to show where the warning was created)
{"level":"INFO","timestamp":"2025-03-25T04:43:39.251Z","logger":"kafkajs","message":"[Consumer] group: 'order-group' subscribed to topic: 'orders-topic'"}
{"level":"ERROR","timestamp":"2025-03-25T04:43:39.298Z","logger":"kafkajs","message":"[Consumer] group: 'order-group' failed to subscribe to topic: 'orders-topic'"}
{"level":"INFO","timestamp":"2025-03-25T04:43:39.641Z","logger":"kafkajs","message":"[Consumer] group: 'order-group' subscribed to topic: 'orders-topic'"}
{"level":"INFO","timestamp":"2025-03-25T04:43:43.547Z","logger":"kafkajs","message":"[Consumer] group: 'order-group' subscribed to topic: 'orders-topic'"}
{"topic":"orders-topic",
  partition: 1,
  offset: '0',
  key: undefined,
  value: {
    orderId: 'order1',
    orderAmount: 100,
    status: 'PENDING',
    timestamp: '2025-03-25T04:44:50.090Z'
  }
}
```

5. Develop a Payment Consumer (paymentConsumer.js)

Consume messages from payments-topic

Log the following clearly:

- o Partition number
- o Offset
- o Parsed JSON data fields (paymentId, orderId, etc.)

```
const { Kafka } = require('kafkajs')

const kafka = new Kafka({
  clientId: 'payment-consumer',
  brokers: ['localhost:9092']
})

const consumer = kafka.consumer({ groupId: 'payment-group' })

const consume = async () => {
  await consumer.connect()
  await consumer.subscribe({ topic: 'payments-topic', fromBeginning: true })

  await consumer.run({
    eachMessage: async ({ topic, partition, message }) => {
      const parsedMessage = JSON.parse(message.value.toString())
      console.log({
        topic: topic,
        partition: partition,
        offset: message.offset,
        key: message?.key?.toString(),
        value: parsedMessage,
      })
    },
  })
}

consume()
```

```
○ (base) spartan@MLK-SCS-P0WGL9N2QF HW8 % node paymentConsumer.js
(node:54280) TimeoutNegativeWarning: -1742877858426 is a negative number.
Timeout duration was set to 1.
(Use `node --trace-warnings ...` to show where the warning was created)
{"level":"INFO","timestamp":"2025-03-25T04:44:18.437Z","logger":"kafkajs","message":{}
{"level":"INFO","timestamp":"2025-03-25T04:44:21.516Z","logger":"kafkajs","message":{}
{"payments-topic": [0,1,2], "groupProtocol": "RoundRobinAssigner", "duration": 3077}
{
  topic: 'payments-topic',
  partition: 1,
  offset: '0',
  key: undefined,
  value: {
    paymentId: 'payment1',
    orderId: 'order1',
    paymentAmount: 50,
    status: 'COMPLETED',
    timestamp: '2025-03-25T04:45:01.343Z'
  }
}
```